

# Dynamic Programming

## Chapter 0: Preface and Overview

Thomas J. Sargent and John Stachurski

2023

# What is this book about?

These are the companion slides to [Dynamic Programming \(Vol 1\)](#)

Dynamic programming has a **vast** array of applications

- robotics
- artificial intelligence
- computational biology
- management science
- finance

Used daily to

- sequence DNA
- manage inventories
- control aircraft
- route shipping
- test products
- optimize database operations
- recommend products
- solve research problems
- etc., etc.

**Example.** Nvidia Hopper GPU DPX instruction sets are aimed specifically at accelerated dynamic programming

Within economics and finance, dynamic programming is applied to

- unemployment and search
- monetary policy and fiscal policy
- asset pricing and portfolio choice
- firm investment
- firm entry and exit
- wealth dynamics
- commodity pricing
- sovereign default
- economic geography
- international trade
- dynamic pricing, etc., etc., etc., etc.

# What's covered in the book?

Standard dynamic programming topics, such as

- Markov decision processes
- Optimal stopping

Modern / advanced topics such as

- state-dependent discounting
- recursive preferences
- ambiguity / robust control
- transformations ( $Q$ -transforms, expected value functions, etc.)
- continuous time
- etc.

# Common Symbols

$[m]$	the integers $\{1, \dots, m\}$
$\mathbb{1}\{P\}$	equals 1 if statement $P$ true, 0 otherwise
$\alpha := 1$	$\alpha$ is defined as equal to 1
$\alpha \equiv 1$	function $\alpha$ is everywhere equal to 1
$\mathbb{N}, \mathbb{Z}$ and $\mathbb{R}$	natural numbers, integers and real numbers
$\mathbb{C}$	complex numbers
$\mathbb{Z}_+, \mathbb{R}_+, \text{etc.}$	the nonnegative elements of $\mathbb{Z}, \mathbb{R}, \text{etc.}$
$ x $ for $x \in \mathbb{R}$	absolute value of $x$
$ \lambda $ for $\lambda \in \mathbb{C}$	modulus of $\lambda$
$a \vee b$	$\max\{a, b\}$
$a \wedge b$	$\min\{a, b\}$

# Common Symbols

$\mathbb{1}$	a function (or vector) everywhere equal to 1
$ B $	the cardinality of set $B$
$\mathbb{R}^n$	all $n$ -tuples of real numbers
$x \leqslant y \ (x, y \in \mathbb{R}^n)$	$x_i \leqslant y_i$ for $i = 1, \dots, n$ (pointwise partial order)
$x \ll y \ (x, y \in \mathbb{R}^n)$	$x_i < y_i$ for $i = 1, \dots, n$
$\mathcal{D}(F)$	the set of distributions (or PMFs) on set $F$
$\mathbb{R}^X$	all functions from $X$ to $\mathbb{R}$
$\mathcal{L}(\mathbb{R}^X)$	the set of linear operators from $\mathbb{R}^X$ to itself
$\mathcal{M}(\mathbb{R}^X)$	all $P \in \mathcal{L}(\mathbb{R}^X)$ with $P \geqslant 0$ and $P\mathbb{1} = \mathbb{1}$

# Julia language: two minute introduction

We use Julia for embedded code because Julia is

- open source
- modern and well-designed
- efficient

Moreover Julia code can be close to underlying equations

- Makes it easy to write and debug



All source code can be found at

<https://github.com/QuantEcon/book-dp1/>

Note that

- Python code is also available
- Julia code is written for clarity, not speed

# Julia Syntax

- Install from <https://julialang.org/> (if you wish)
- To import Library, write **using** Library
- $f(x) = 2x$  defines the function  $f(x) = 2x$
- `cos.(x)` applies `cos` to each elements of vector `x`
- `x.^2` squares each element of vector `x`
- loops / conditions very similar to Python
- data is pass by reference
- some nice multiple dispatch operations

---

## # Defining functions, using conditions and loops

```
function f(x, y)                                # define a function
    if x < y                                     # branch
        return sin(x + y)
    else
        return cos(x + y)
    end
end

function print_plurals(list_of_words)          # define a function
    for word in list_of_words                  # loop
        println(word * "s")
    end
end
```

---

**using** LinearAlgebra

*# import LinearAlgebra library*

$f(x) = 2x$

*# simple function definition*

$f(x) = \text{norm}(x)$

*# norm defined in LinearAlgebra*

$g(x) = \text{sum}(x + x.^2)$

*# dot for pointwise operations*

$\alpha, \beta = 2.0, -2.0$

*# unicode symbols*

$q(x) = \sin(\cos(x))$

*# another function*

$\text{println}(q(5))$

*# OK*

$x = \text{rand}(5)$

*# build a random vector*

$\text{println}(q.(x))$

*# OK*

$\text{println}(q(x))$

*# error*

---

The **composition** of  $f: A \rightarrow B$  and  $g: B \rightarrow C$  is  $g \circ f: A \rightarrow C$  defined by

$$a \mapsto g(f(a))$$

**Example.**  $f(x) = x \wedge 0$  and  $g(x) = x \vee 0$  implies  $g \circ f \equiv 0$

In Julia we can compose as follows

---

```
f(x) = min(x, 0)
g(x) = max(x, 0)
h = f ∘ g           # type \circ and then hit tab
```

---

Further information:

- QuantEcon resources: <https://julia.quantecon.org>
- Other resources: <https://julialang.org/learning/>

## For advanced readers

For the rest of these chapter slides we discuss more advanced topics

- suitable for readers who already know some dynamic programming
- give more idea of what we cover in the book

Less advanced readers can safely skip to Chapter 1

# Optimal consumption

Some readers will have studied an optimization problems such as

$$\max_{(C_t)_{t \geq 0}} \sum_{t \geq 0} \beta^t u(C_t)$$

subject to

$$W_{t+1} = R(W_t - C_t) \quad \text{and} \quad W_t, C_t \geq 0$$

- $W_t$  is current wealth
- $C_t$  is current consumption
- $u(C_t)$  is current utility (reward)
- $\beta$  is a discount factor (time preference)



They will know how to set up the **Bellman operator**

$$(Tv)(w) = \max_{0 \leq c \leq w} \{u(c) + \beta v(R(w - c))\}$$

They will know that, under some conditions,

1.  $T$  is a contraction mapping
2. the unique fixed point of  $T$  is the value function  $v^*$
3.  $v^*$  can be approximated via  $v^* = \lim_{k \rightarrow \infty} T^k v$  for some  $v$
4. optimal consumption at wealth  $w$  can be found by solving

$$c^* \in \operatorname{argmax}_{0 \leq c \leq w} \{u(c) + \beta v^*(R(w - c))\}$$

But is this the best way?

Perhaps we should be using

- time iteration
- Howard policy iteration
- optimistic policy iteration

What are these algorithms?

Do they always converge?

Are they faster / more accurate?

Moreover, isn't our model too simplistic?

Possible extensions / modifications include

- labor income, work-leisure choice, multiple assets
- state-dependent discounting / preference shocks
- Epstein–Zin preferences
- ambiguity and ambiguity aversion
- robustness / risk-sensitivity / adversarial agents
- expected value formulation
- quantile preferences
- continuous time
- etc.

**Example.** What if discounting depends on the action (e.g., Uzawa preferences, as in, say, Schmitt-Grohé Uribe 2004)?

- $\beta$  depends on consumption (and maybe labor/leisure)

$$(Tv)(w) = \max_{0 \leq c \leq w} \{u(c) + \beta(c)v(R(w - c))\}$$

- Is  $T$  still a contraction?
- Are all the previous optimality results still valid?
- Which algorithms converge?

**Example.** What if we face **state-dependent discounting** (as in, say, Krusell Smith 1998)?

- $\beta$  now depends on an exogenous state process

$$(Tv)(w, z) = \max_{0 \leq c \leq w} \left\{ u(c) + \beta(z) \sum_{z'} v(R(w - c), z') Q(z, z') \right\}$$

- Is  $T$  still a contraction?
- Are all the previous optimality results still valid?
- Which algorithms converge?

**Example.** What about state-dependent discounting via **expected value functions** (as in the structural estimation literature), where

$$g(w, z, c) := \sum_{z'} v(R(w - c), z') Q(z, z')$$

with “Bellman operator”

$$(Hg)(w, z, c) =$$

$$\sum_{z'} \max_{0 \leq c' \leq R(w - c)} \{u(c') + \beta(z') g(R(w - c), z', c')\} Q(z, z')$$

Does  $H$  have all the same properties as  $T$ ?

What are the equivalent algorithms and do they converge?

And what happens if we introduce **Epstein–Zin preferences** (e.g., Bansal Yaron 2004, Basu Bundick 2017)?

$$(Tv)(w, z) =$$

$$\max_{0 \leq c \leq w} \left\{ c^\alpha + \beta(z) \left[ \sum_{z'} v(R(w - c), z')^\gamma Q(z, z') \right]^{\alpha/\gamma} \right\}^{1/\alpha}$$

- Is  $T$  still a contraction?
- Are all the previous optimality results still valid?
- Which algorithms converge?

Or **risk-sensitive preferences** (as in, say, Hansen Sargent 2007)?

$$(Tv)(w, z) =$$

$$\max_{0 \leq c \leq w} \left\{ u(c) + \frac{\beta(z)}{\theta} \ln \left[ \sum_{z'} e^{\theta v(R(w-c), z')} Q(z, z') \right] \right\}$$

- Is  $T$  still a contraction?
- Are all the previous optimality results still valid?
- Which algorithms converge?



And what happens if we introduce

- quantile preferences?
- adversarial agents?
- ambiguity?

Is the Bellman operator a contraction?

Do the optimality properties hold?

Which algorithms converge?

How do we compute solutions?

We will address these questions by

1. covering the basic, foundational models
2. introducing state-dependent discounting
3. introducing recursive preferences
4. developing a general framework to handle all of the above
5. studying optimality and algorithms in the general framework
6. studying relationships between dynamic programs
7. switching to continuous time

A quick sketch of the main ideas:

Our general analysis uses **abstract Bellman equations** of the form

$$v(x) = \max_{a \in \Gamma(x)} B(x, a, v)$$

Includes all models discussed so far as special cases

**Example.** For the Epstein–Zin specification on slide 23, we use

$$B(x, a, v) = B((w, z), c, v) =$$

$$\left\{ c^\alpha + \beta(z) \left[ \sum_{z'} v(R(w - c), z')^\gamma Q(z, z') \right]^{\alpha/\gamma} \right\}^{1/\alpha}$$

Behavior is determined via policy functions

A **feasible policy** is a map  $\sigma: X \rightarrow A$  such that

$$\sigma(x) \in \Gamma(x) \text{ for all } x \in X$$

- always respond to state  $x$  with action  $x := \sigma(x)$

Let  $\Sigma$  = the set of all feasible policies

Given  $\sigma \in \Sigma$ , suppose  $v_\sigma$  satisfies

$$v_\sigma(x) = B(x, \sigma(x), v_\sigma) \quad \text{for all } x \in X$$

Interpretation:

- $v_\sigma(x) =$  **lifetime value** of policy  $\sigma$  given  $X_0 = x$

Questions

- Is this interpretation reasonable?
- When is  $v_\sigma$  well defined / uniquely defined?

Suppose  $v_\sigma$  is uniquely defined for all  $\sigma \in \Sigma$

Then we can introduce the **value function** via

$$v^*(x) := \max_{\sigma \in \Sigma} v_\sigma(x)$$

= max lifetime value from state  $x$

A policy  $\sigma \in \Sigma$  is called **optimal** if

$$v_\sigma = v^*$$

Under what conditions do standard optimality results hold?

- does  $v^*$  satisfy  $v^*(x) = \max_{a \in \Gamma(x)} B(x, a, v^*)$ ?
- does an optimal policy exist?
- are optimal policies characterized by Bellman's principle, where

$$\sigma(x) \in \operatorname{argmax}_{a \in \Gamma(x)} B(x, a, v^*)$$

- under what conditions on  $B$  do the usual algorithms converge?
- what transformations can we apply to  $B$  in order to simplify analysis / computation?



We address these questions step by step

1. Provide general conditions
2. Provide more specialized conditions for special cases
  - contractive models
  - eventually contractive models
  - concave models / convex models
3. Connect these conditions to applications

We also discuss “completely abstract” DP models that can handle additional applications

- See Ch 9

# Advantages

Working in this abstract setting

- simplifies proofs
- clarifies theory
- clarifies relationships between “similar” dynamic programs

Notice

- no explicit dynamics
- no measure theory

Such details are pushed back to applications

- are the conditions we place on  $B$  satisfied?

# Qualifications

Most of the current volume focuses on finite states and actions

This is not a bad thing because

- eliminates lots of qualifying statements
- minimizes functional analysis / measure theory
- covers all computable models

For general state spaces, see

1. the abstract theory in Ch 9 and
2. Vol II (forthcoming!)

# Structure

The book starts with concrete problems

1. finite horizon search
2. infinite horizon search
3. Markov dynamics
4. optimal stopping
5. Markov decision processes
6. state-dependent Markov decision processes

Then we shift to abstract problems (Ch.s 7–9)